

F5 NGINX


Lessons Learned: Adding OpenTelemetry to a Modern App

Dave McAllister

Every company is on a cloud journey

To increase velocity, agility and responsiveness



Retain & Optimize



Tightly Coupled Apps,
Slow Deployment Cycles



Lift & Shift




Primarily using
Cloud IaaS




Re-Factor




Cloud Managed e.g. RDS,
DynamoDB, SaaS




More Modular, but
Dependent App Components



Re-Architect / Cloud-Native



Cloud First Architecture




Loosely Coupled Microservices,
and Serverless Functions



“By 2025, 85% of organizations will run containers in production, up from less than 30% in 2020” – Gartner, Dec 14, 2020

Modern Apps



Are defined by their capabilities not their implementation.



We speak about them using terms like portable, scalable, observable, reproducible, and debuggable.

No One Cares if it is all in COBOL - They Care if it Works

When your app can do this, it is modern


- Deploy in minutes
- **Quickly find performance bottlenecks**
- Easily change hosting provider
- Scale up/down quickly
- Gracefully degrade upon failure
- **Provide answers to platform engineers' questions**
- Protect itself from attacks
- Manage state in a knowable way
- **Provide context on errors and crashes**
- Costs scale reasonable with consumption
- Serve the customers' needs

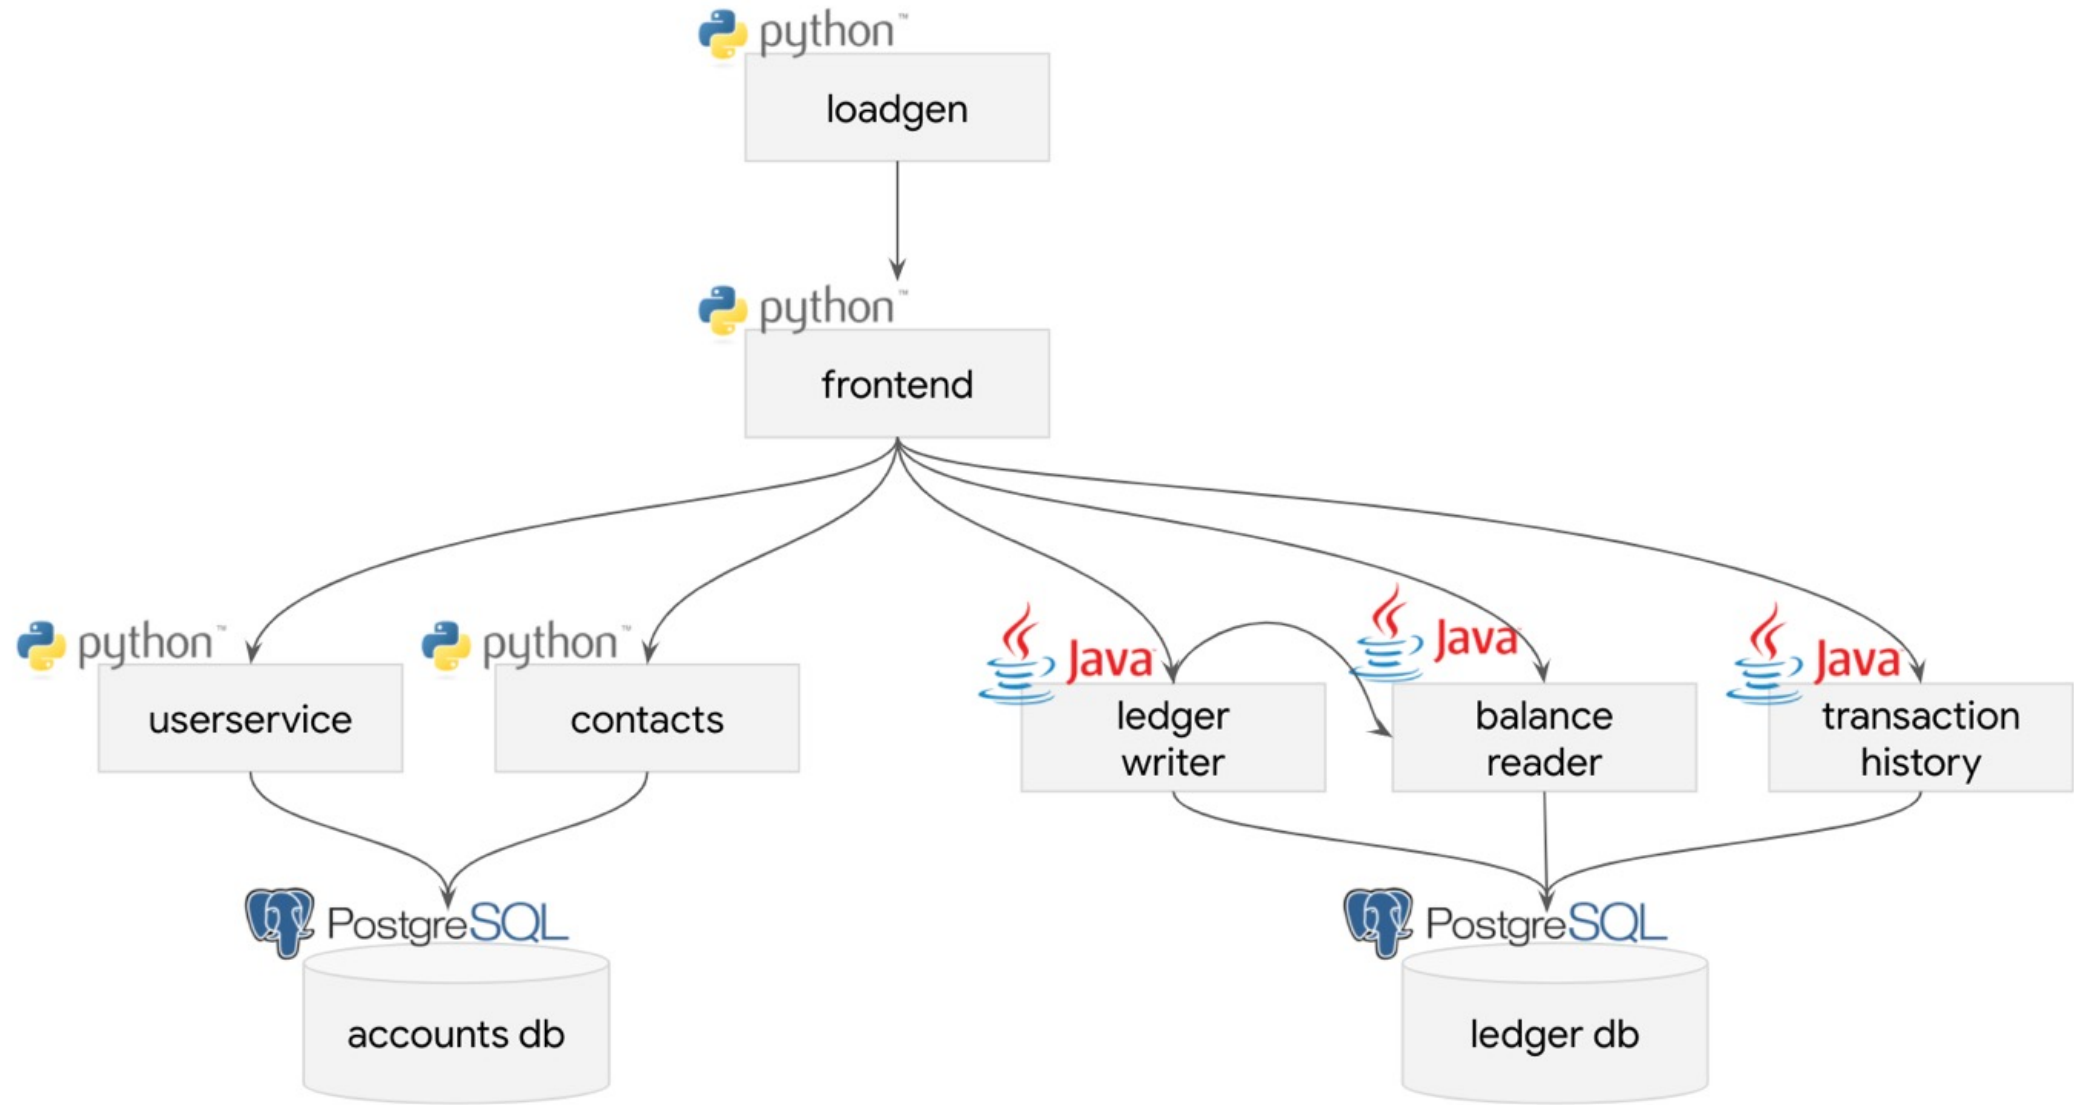
What is MARA?

An microservices architecture using Kubernetes that aims to be as production ready as possible.

Configured and integrated components deployed via Infrastructure as Code

The application is only one part of the modern application deployment

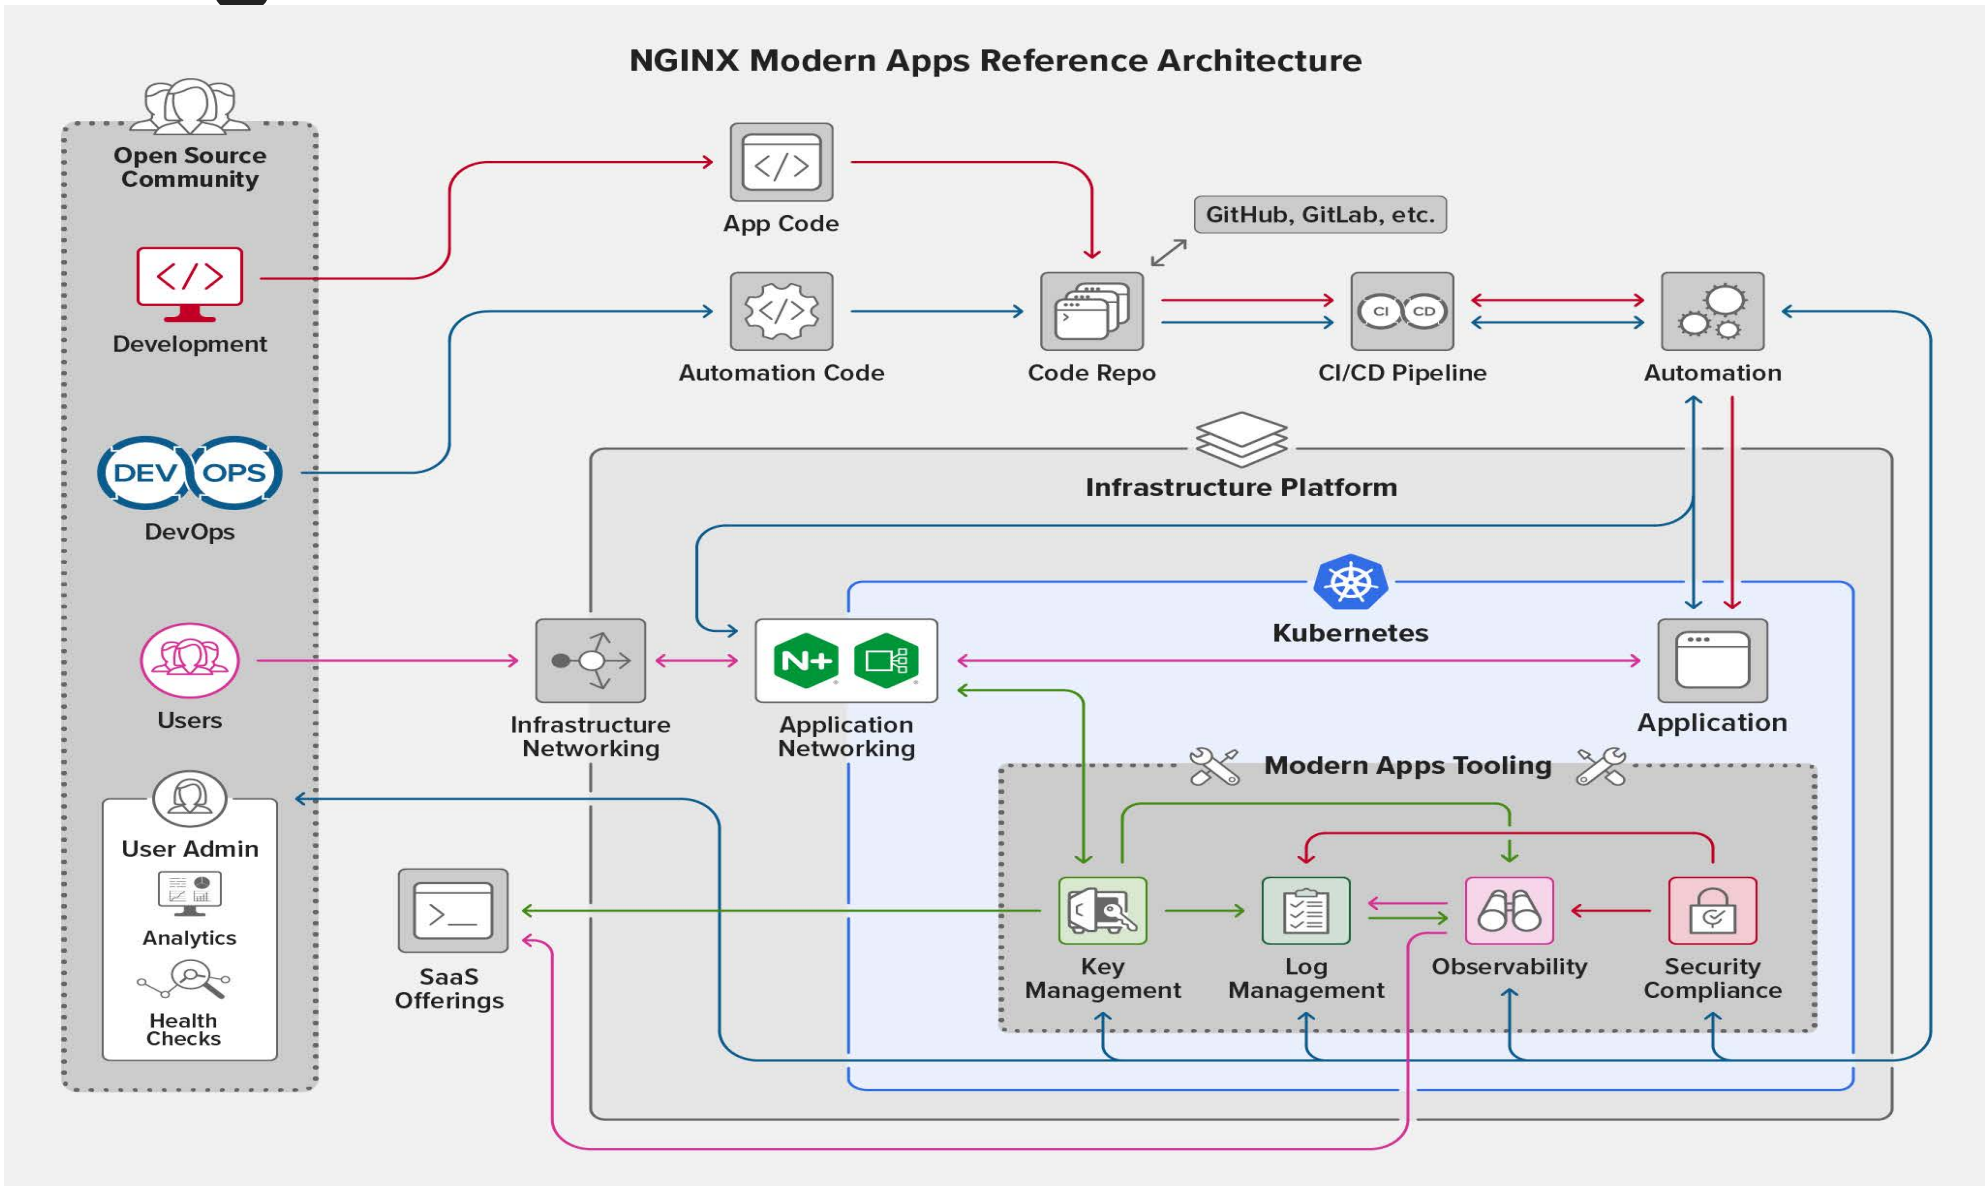




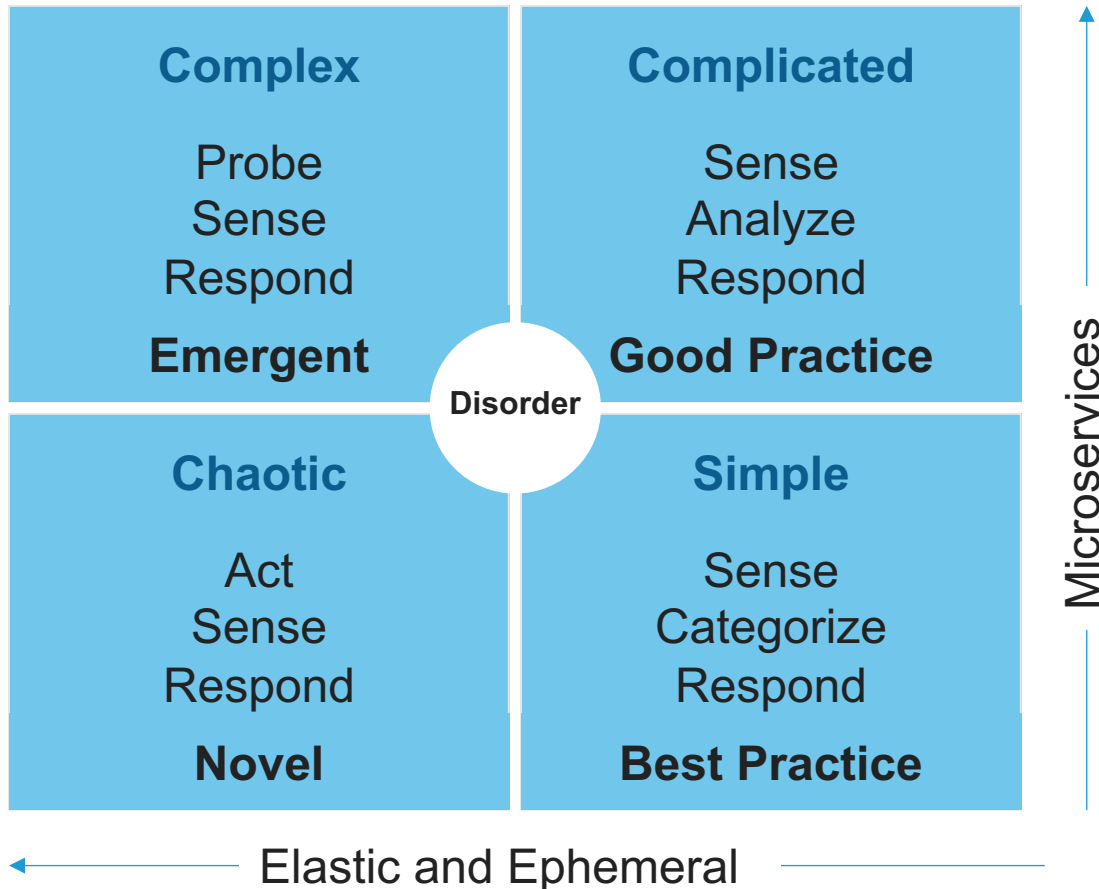
What's in the box?

- Infrastructure: Kubeconfig, AWS, Linode, Digital Ocean
- Observability:
 - Tracing: OpenTelemetry Operator
 - Metrics: Prometheus
 - Logging: Elasticsearch / Filebeat / Kibana
 - Visualization: Grafana
- Management:
 - Certificates: Cert-Manager
 - Secrets Management: Pulumi / Kubernetes
 - Ingress: NGINX Ingress Controller / NGINX Plus Ingress Controller
- Load Generation: Locust
- Application: Bank of Sirius
- Infrastructure as Code: Pulumi in Python
- Continuous Integration: Jenkins

Drawing it all out



Observability answers Challenges



- Microservices create complex interactions.
- Failures don't exactly repeat.
- Debugging multi-tenancy is painful.
- Scale of data is massive
- Faster release cycles
- You own the code

OBSERVABILITY IS A DATA
PROBLEM

THE MORE OBSERVABLE A
SYSTEM, THE QUICKER WE
CAN UNDERSTAND WHY IT'S
ACTING UP AND FIX IT

Full-Stack Visibility & Context-Rich Insights

Metrics

Do I have
a problem?

DETECT

Traces

Where is the
problem?

TROUBLESHOOT

Logs

Why is the
problem
happening?

ROOT CAUSE

Identifying the needs

Logs

Metrics

Distributed
Traces

Error
Aggregation

Health Checks

Heap/Core
Dumps

RT State
Introspection

Wishlist

Technology	Logging	Tracing	Metrics	Error Agg	Health Checks	Runtime Intro	Heap/Core Dumps
Elastic APM	Yes	Yes	Yes	Yes	Yes	No	No
Grafana	Yes	Yes	Yes	Yes	Yes	No	No
Graylog	Yes	No	No	No	No	No	No
Jaeger	No	Yes	No	Yes	No	No	No
OpenCensus	No	Yes	Yes	No	No	No	No
OTel	Beta	Yes	Yes	Yes	Yes	No	No
Prometheus	No	No	Yes	No	No	No	No
Statsd	No	No	Yes	No	No	No	No
Zipkin	No	Yes	No	Yes	No	No	No

Time to get Qualitative

OpenCensus and OpenTracing have merged into OpenTelemetry!




OpenCensus

Easily collect telemetry like metrics and distributed traces from your services

The OpenTracing project is *archived*. [Learn more.](#)
[Migrate to OpenTelemetry](#) today!

What is OpenTelemetry?

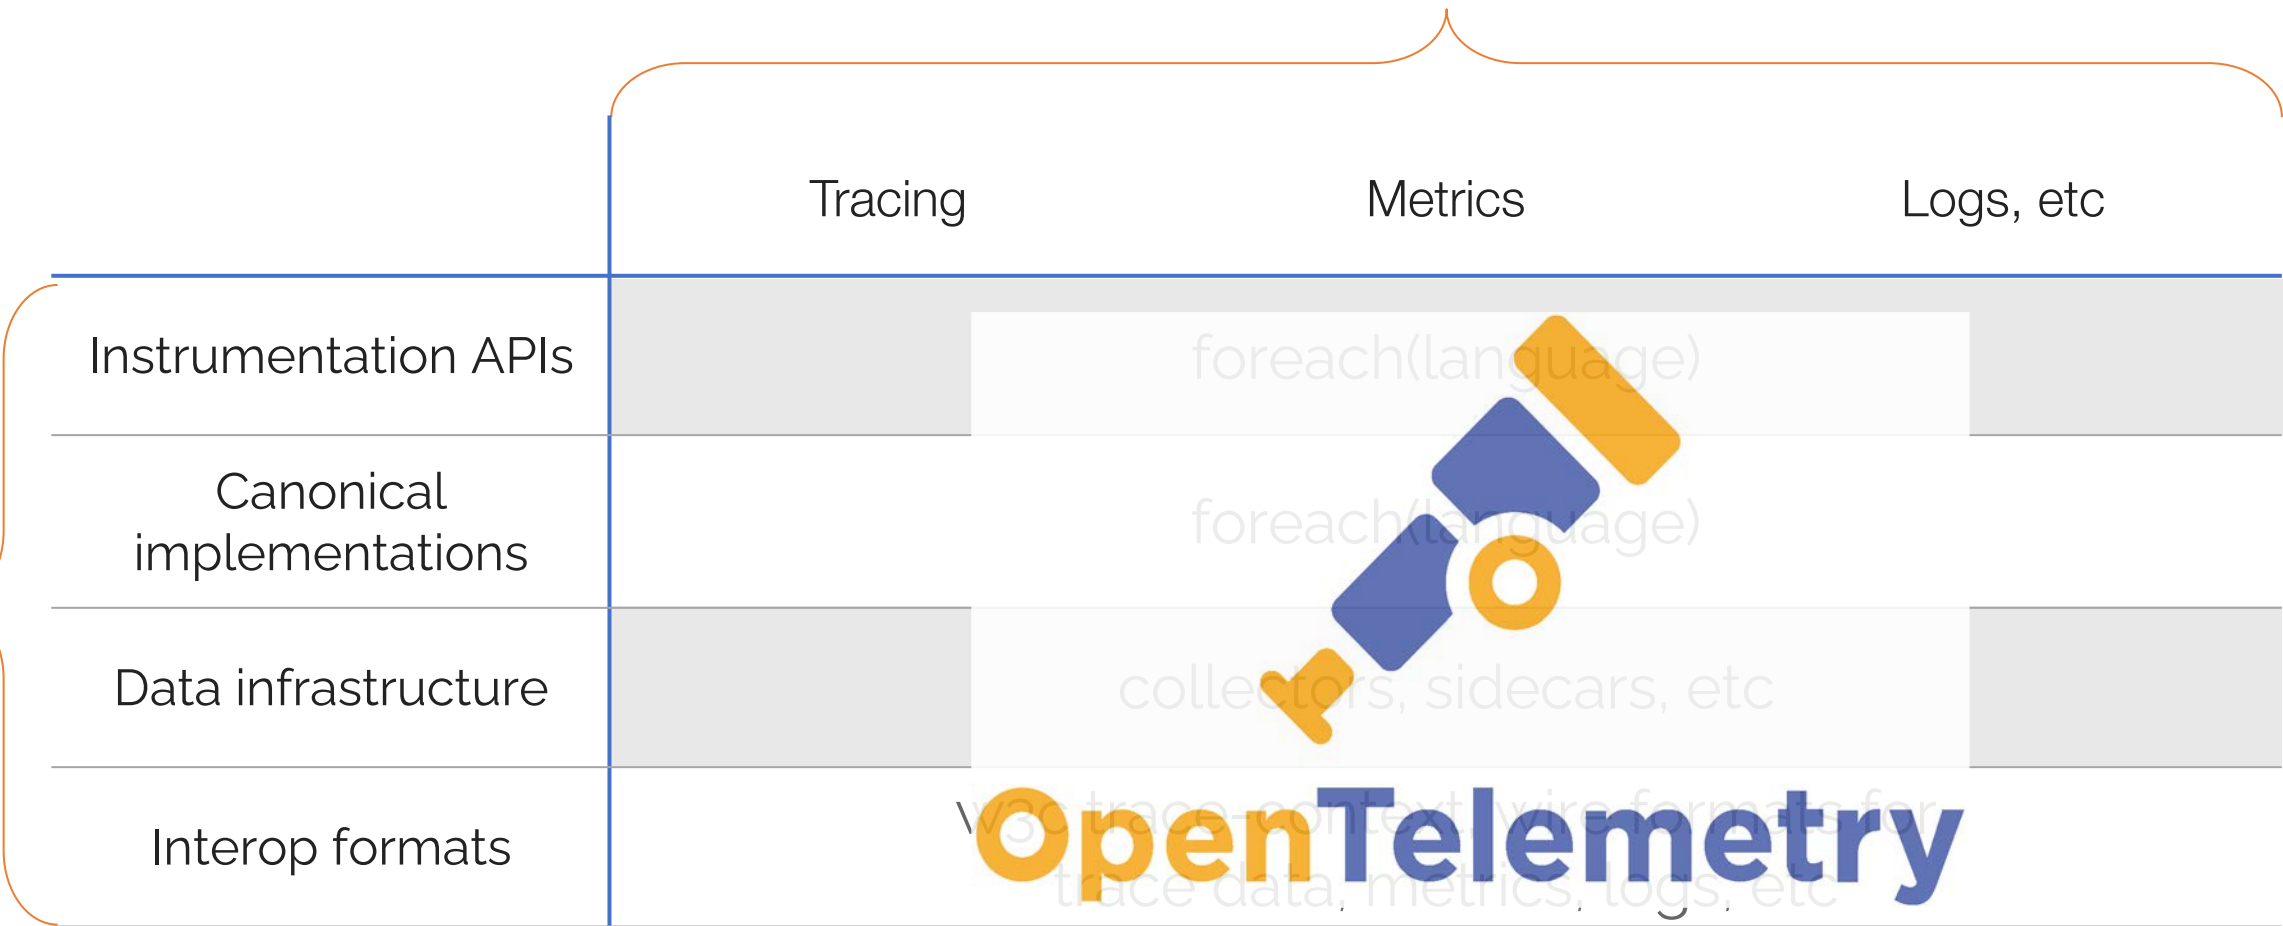
- Standards-based agents, cloud-integration
- Automated code instrumentation
- Support for developer frameworks
- Any code, any time



Cloud Native Telemetry

Telemetry “verticals”

Telemetry
“layers”



Let's start with logs


We all know what logs are but their simplicity rapidly leads to some complex decisions

- Simple - Harvest the log output
- Complicated – Transport, Storage, indexing (?), lifetime

To be useful, our log files must be easily searchable based on varying criteria

Elastic Stack, for now

- Filebeat – data transport in Kubernetes DaemonSet
- Bitnami – chart to split the deployment into ingest, co-ordinating, master and data nodes
- Kibana – search + pre-loaded stuff



It works, but


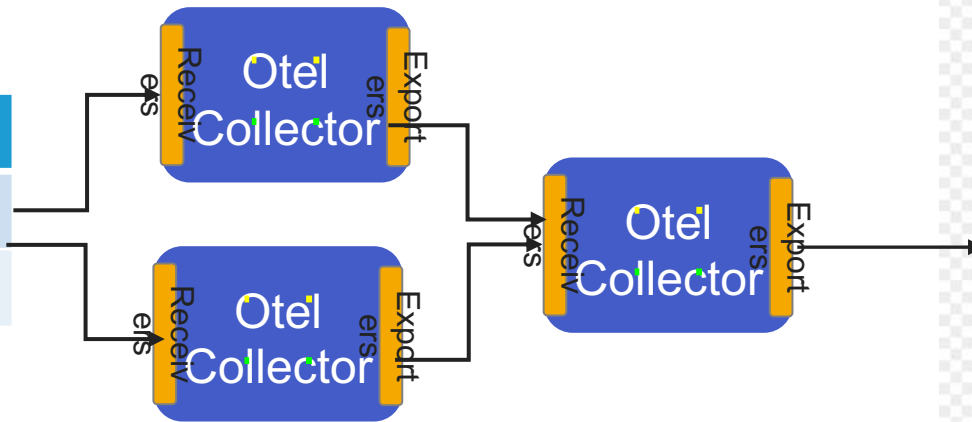
- Extremely resource hungry
- Query variance is okay

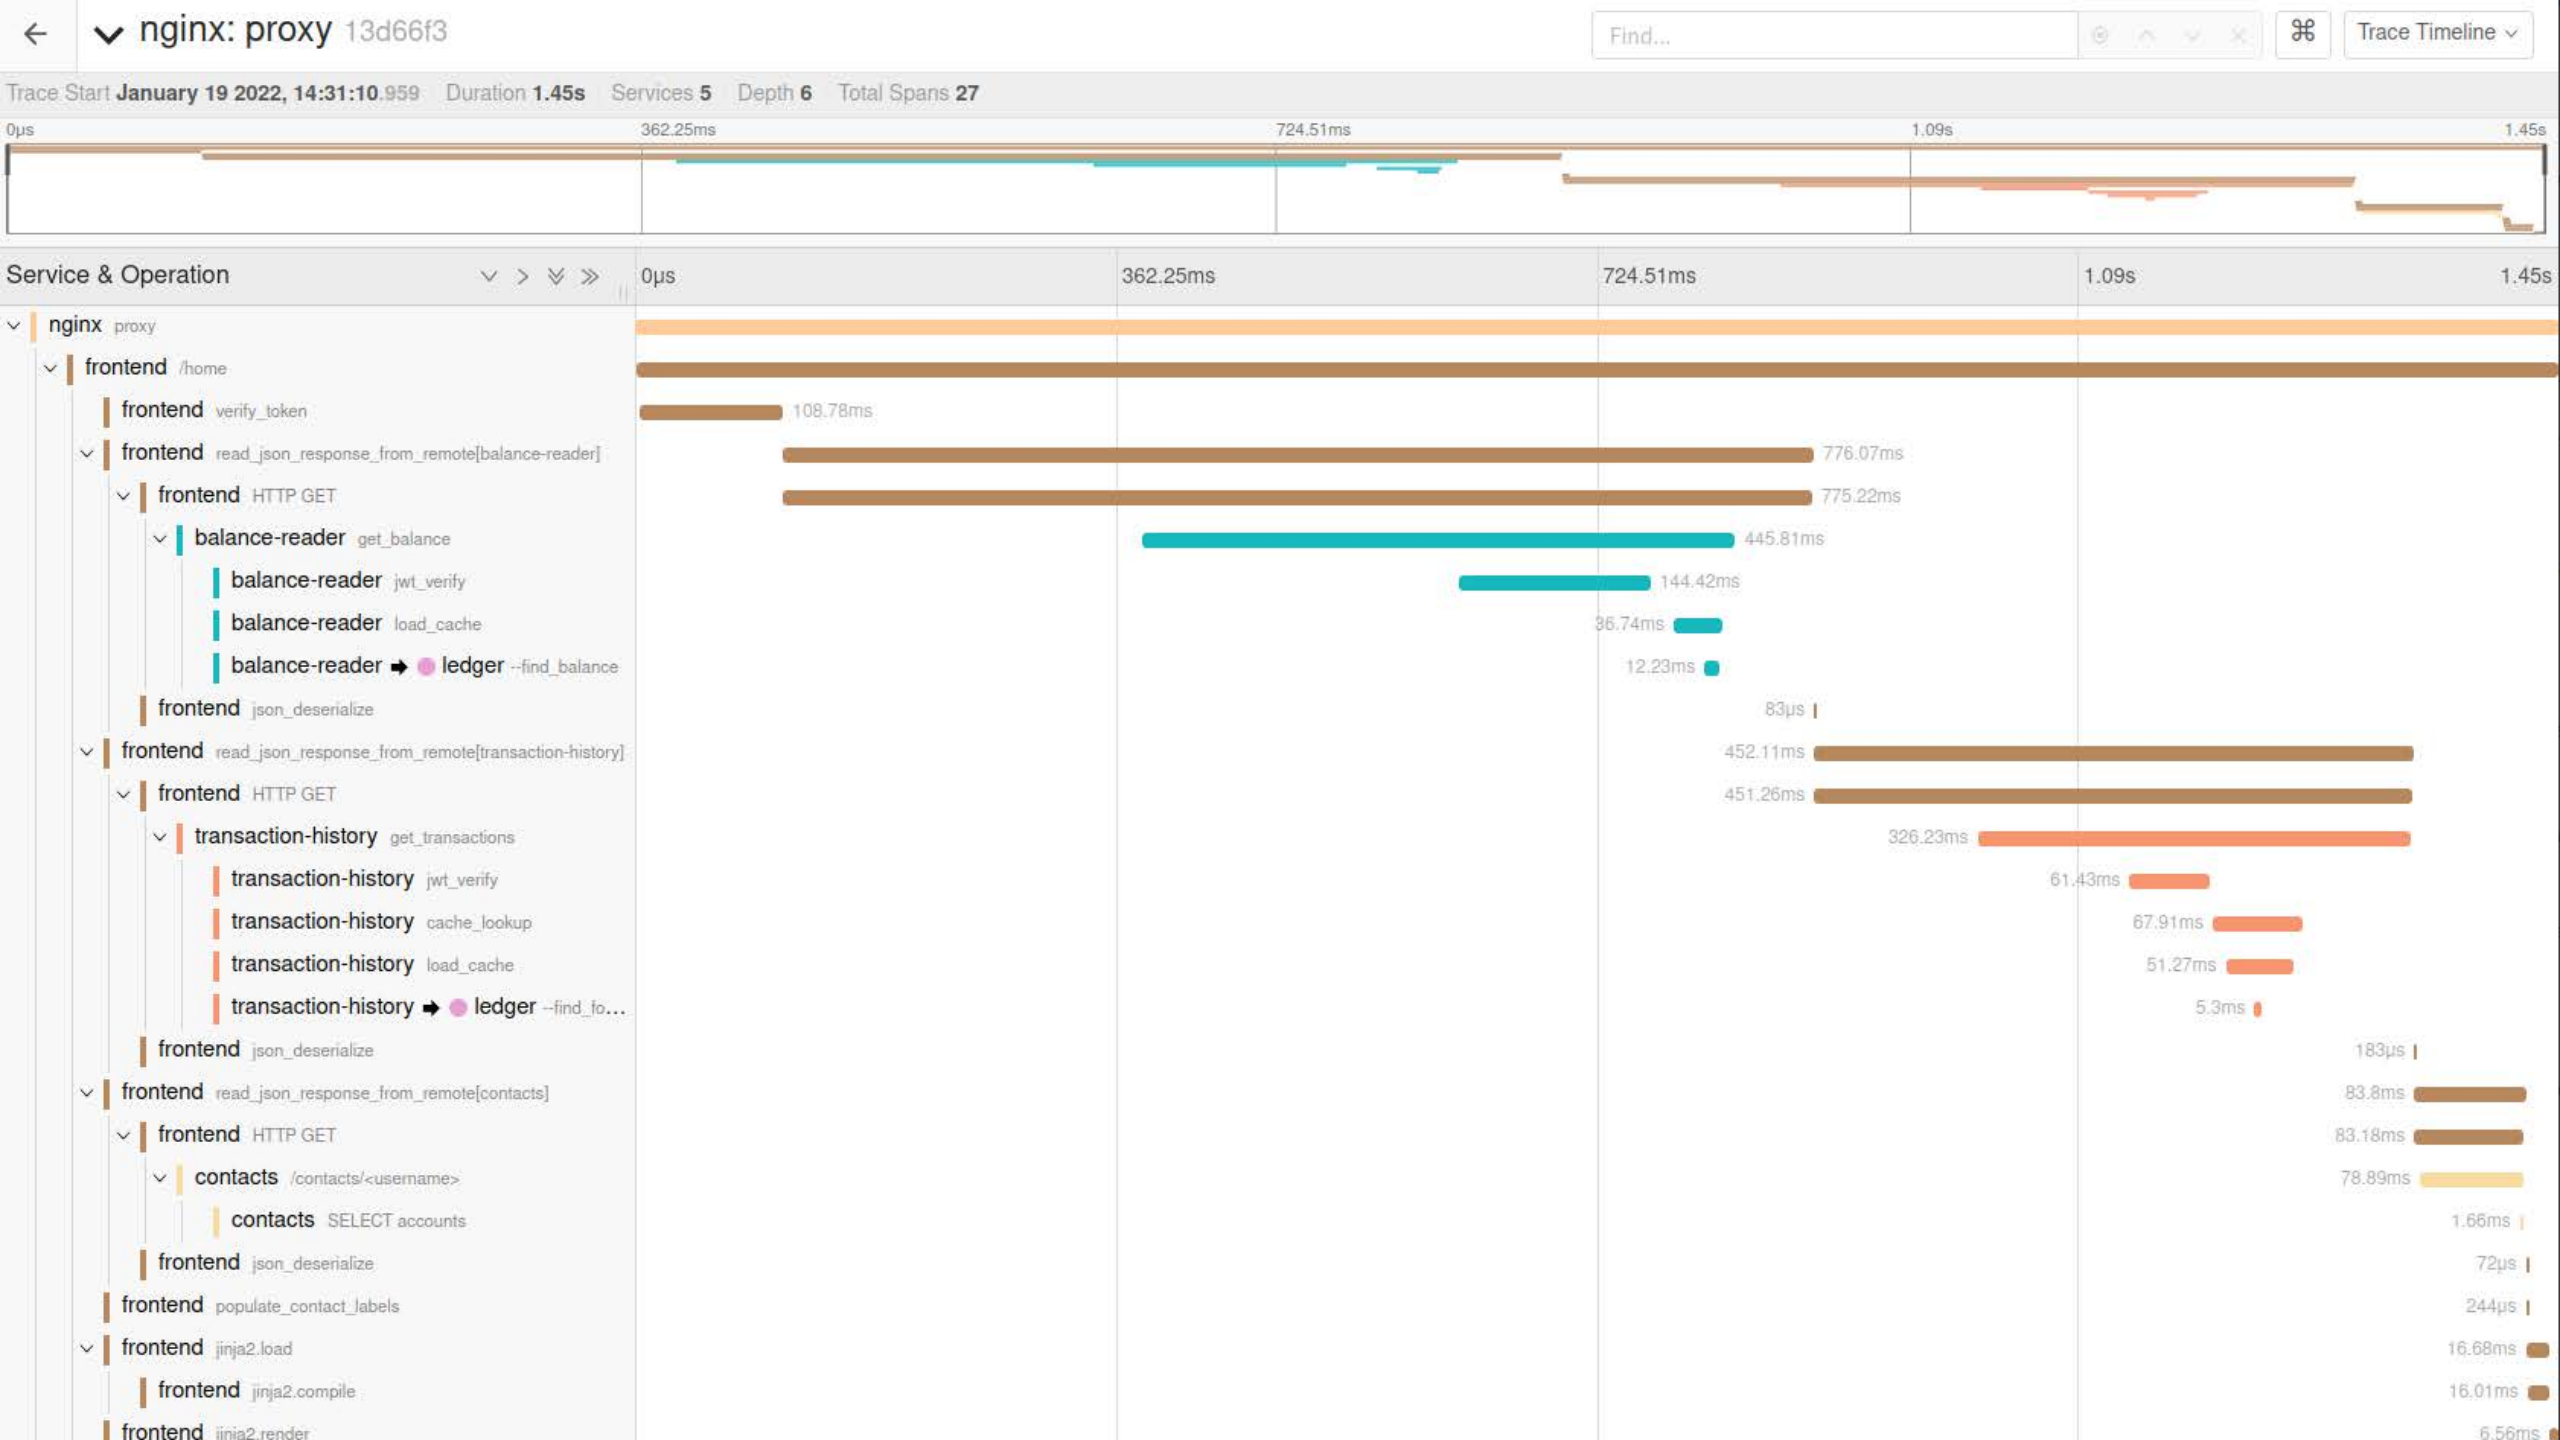
Distributed Tracing

Complex and semi-chaotic

- Must not impact QoS of the application
- Support all desired languages

Language	Framework	Services
Java	Sprint Boot	3
Python	Flask	3






But it's all about the language


Python

- Pretty straightforward
 - Added 2 files
 - Updated requirements.txt to include dependencies




dekobon feat: use bunyan output for json logs with python apps ...

..




Dockerfile

feat: add trace ids to Python logs



gunicorn.conf.py

feat: use bunyan output for json logs with python apps



tracing.py

feat: add trace ids to Python logs

And some were easier than others

Java

- Straight Java / Greenfield = Not too bad
 - Import the libraries and use the APIs
- With Spring, life looked easy
 - Use Spring Cloud Sleuth
 - Adds trace/span IDs to Slf4j
 - Instruments common ingress and egress points
 - Adds traces to scheduled tasks
 - Can generate Zipkin traces
- But at that time:
 - Autoconfig was a milestone release, and supported old Otel versions
 - We needed to pull from Spring Snapshot due to coded dependency references

Answer: a common telemetry module

- Extended tracing functionality
 - Provided Spring enabled autoconfiguration classes, adding more trace resource attributes
 - Built a NoOp implementation to let us disable tracing
 - Added a trace name interceptor to standardize our trace names
 - Added an error handler to output errors both to logs and traces
 - Enhanced the implementation of tracing attributes (service name, instance id, machine id, etc)
 - Built a tracing statement inspector to put trace ids into comments that precede SQL statements

We also extended the reach to Apache by creating a Spring compatible HTTP client

All this was integrated using the OpenTelemetry NGINX module (beta)

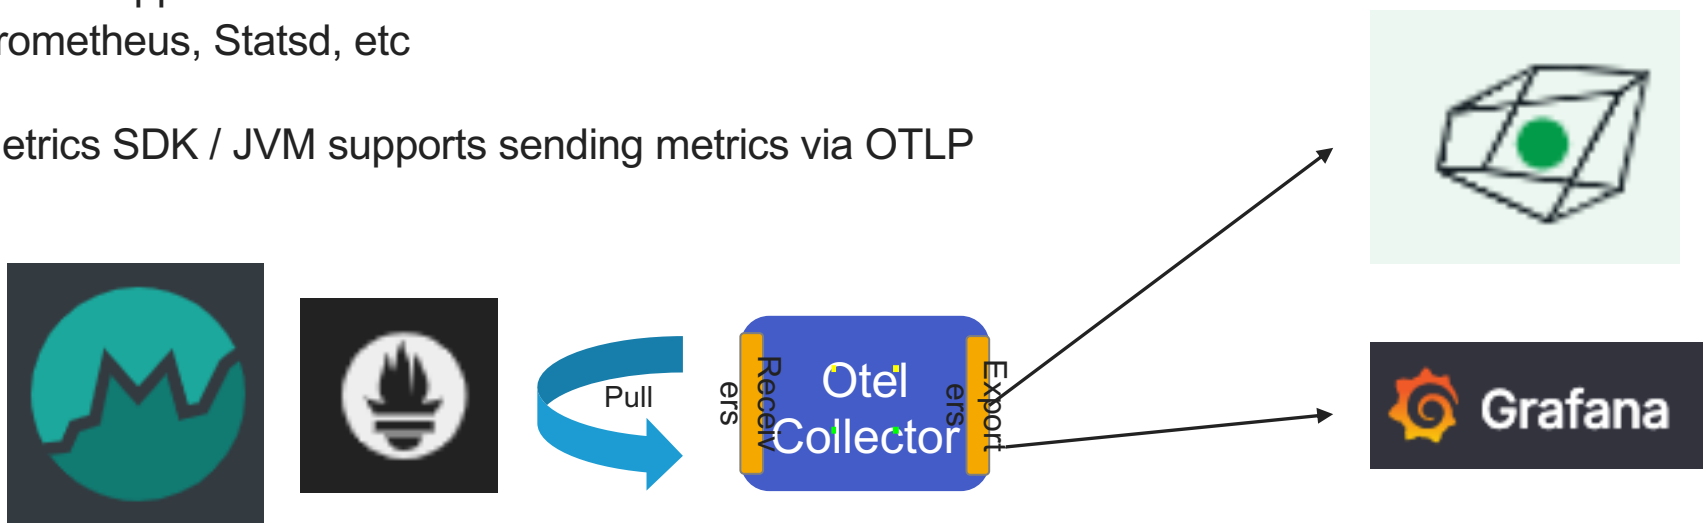
[nginxinc/nginx-unsupported-modules: Container builds of unsupported NGINX modules \(github.com\)](https://github.com/nginxinc/nginx-unsupported-modules)

Counting on Metrics

- We skipped Python
- Java required some study and decisions.
 - Original code (Bank of Anthos) used Micrometer/Stackdriver
- We found there were some significant limits to Otel for metrics in JVM
- Micrometer is a mature metrics layer for JVM
 - It is the default metrics API in Spring

OpenTelemetry and Micrometer

- OTEL Collector can use metrics from *anything*
 - Prometheus, Statsd, OTLP, etc
- Micrometer supports a lot of back ends
 - Prometheus, Statsd, etc
- OTEL Metrics SDK / JVM supports sending metrics via OTLP



Summary

Otel FTW

Distributed Tracing

- Java: Spring Cloud Sleuth → Spring Cloud Sleuth OTEL exporter → OTEL Collector → Pluggable Store
- Python: OTEL Python libraries → OTEL Collector → Pluggable Store
- NGINX (Ingress Controller is not supported yet): NGINX OTEL module → OTEL Collector → Pluggable Store

Metrics

- Java: Micrometer via Spring → Prometheus exporter → OTEL Collector
- Python WSGI: GUnicorn statsd → Prometheus (via statsd / ServiceMonitor)
- Python: No implementation yet
- NGINX: Prometheus Endpoint → Prometheus (via ServiceMonitor)

Logs

- All Container Logs: Filebeat → Elasticsearch / Kibana

Summary

Error Aggregation

- OTEL distributed traces → Pluggable Store search

Health Checks

- Java: Spring Boot Actuator → Kubernetes
- Python: Flask Management Endpoints Python module → Kubernetes

Runtime Introspection

- Java: Spring Boot Actuator
- Python: Flask Management Endpoints Python module

Heap/Core Dumps

- Java: Spring Boot Actuator support for thread dumps
- Python: no support yet

TL;DR

Metrics, Traces and Logs

- All took different approaches to get what we wanted
- The Collector was our friend

Metrics and Traces had some interesting gotchas

This is a snapshot in time; things have changed

Auto-configuration is great

- When it works
- And gives you what you want



<https://github.com/nginxinc/kic-reference-architectures>

Try it for yourself



[nginxinc/bank-of-sirius: Bank of Sirius \(github.com\)](https://github.com/nginxinc/bank-of-sirius)

Questions

Join our Open Source
Community slack!

