

How to achieve the scalability, high availability, and elastic ability of your database infrastructure on Kubernetes

Trista Pan
panjuan@apache.org

Trista Pan

SphereEx Co-Founder & CTO

Apache Member

AWS Data Hero

Tencent Cloud TVP

Apache ShardingSphere PMC

Apache brpc (Incubating) & Apache AGE

& Apache HugeGraph (Incubating) mentor

China Mulan Community Mentor



Bio: <https://tristazero.github.io>

LinkedIn: <https://www.linkedin.com/in/panjuan>

GitHub: <https://github.com/tristaZero>

Twitter: @tristaZero


Project Twitter: @ShardingSphere

Content


- ✓ SRE & SLA & DBRE
- ✓ The new needs for a database on the cloud
- ✓ Idea & architecture
- ✓ Handling SQL
- ✓ Demo

SRE & SLA & DBRE

- ✓ Database Reliability Engineering (DBRE) is basically a subset of Site Reliability Engineering (SRE)
- ✓ Stateless service VS stateful service (Persistence & status)
- ✓ SLA (Service Level Agreement) & SLO (Service Level Objectives) & SLI (Service Level Indicators)

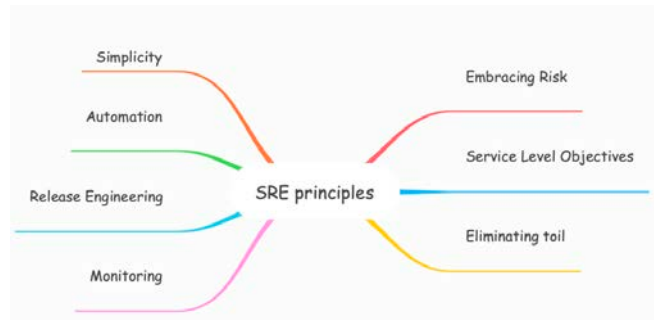


Edited by Betsy Beyer, Chris Jones,
Jennifer Petoff & Niall Richard Murphy



Laine Campbell & Charity Majors

New needs for databases



The needs for a database on the cloud

- ✓ Large data to manage
- ✓ Efficient queries
- ✓ Data security
- ✓ Traffic governance
- ✓ Elastic scaling
- ✓ Backup & recovery
- ✓ Metrics
- ✓ Portability
- ✓ Out-of-the-box deployment



Data Sharding

HA & read/write splitting & traffic strategy

Data Encryption

Monitor

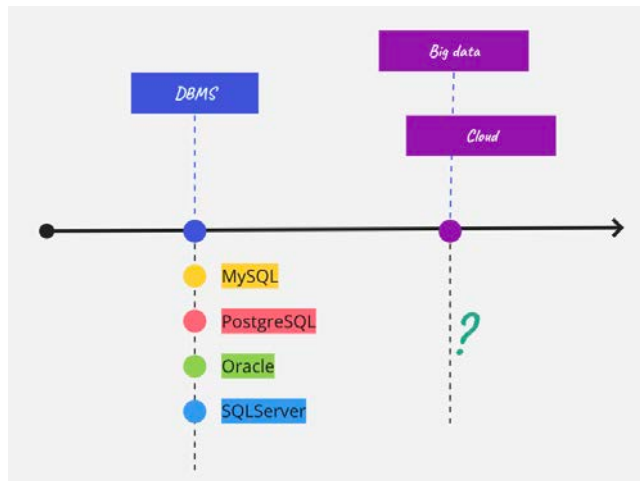
Reshard for computing nodes and storage nodes

Helm & Operator on Kubernetes

Monolithic database on the cloud

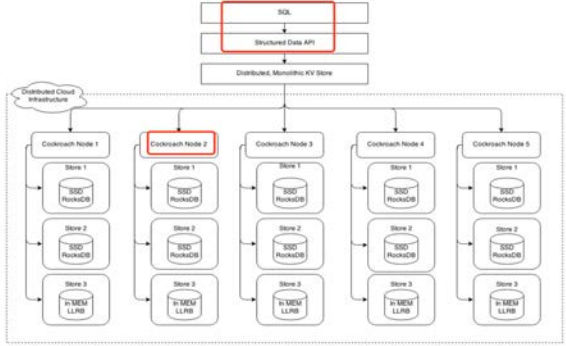
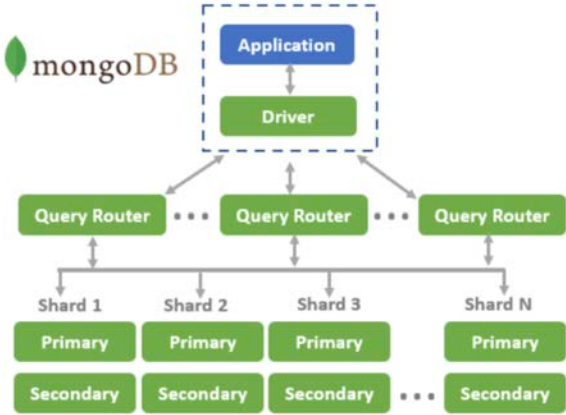



Benefits




- ✓ Leverage the existing databases
- ✓ Upgrade it into a distributed database at low cost
- ✓ SQL audit & Traffic governance & Elastic scaling
- ✓ Solve the headache of moving database into Kubernetes
- ✓ Out-of-the-box deployment
- ✓ No lock-in

Distributed database




Application -> Database



Before

Distribute
database system



After

Apache ShardingSphere

Fork 6k Starred 17.3k

Security Insights

About

Ecosystem to transform any database into a distributed database system, and enhance it with sharding, elastic scaling, encryption features & more

mysql sql database bigdata postgresql shard rdbs distributed-transactions distributed-database dba encrypt database-cluster oltp distributed-sql-database database-plus

Readme

Apache-2.0 license Code of conduct 17.3k stars 1k watching 6k forks

Releases 49

5.2.0 Latest 6 days ago

+ 48 releases

Packages 1

shardingsphere-proxy

Contributors 437



+ 426 contributors

Environments 1

github-pages Active

What is Apache ShardingSphere?

The ecosystem to transform any database into a distributed database system, and enhance it with sharding, elastic scaling, encryption features & more.

Download

Learn More

Academic Publications



ShardingSphere

ShardingSphere English

ShardingSphere > Overview

1. Overview 2. Quick Start 3. Features 4. User Manual 5. Dev Manual 6. Test Manual 7. Reference 8. FAQ 9. Downloads

This chapter mainly introduces what Apache ShardingSphere is, as well as its design philosophy and deployment architecture.

For frequently asked questions, please refer to FAQ.

What is ShardingSphere

Introduction

Apache ShardingSphere is an open source ecosystem that allows you to transform any database into a distributed database system. The project includes a JDBC and a Proxy, and its core adopts a micro-kernel and pluggable architecture. Thanks to its plugin-oriented architecture, features can be flexibly expanded at will.

The project is committed to providing a multi-source heterogeneous, enhanced database platform and further building an ecosystem around the upper layer of the platform. Database Plus, the design philosophy of Apache ShardingSphere, aims at building the standard and ecosystem on the upper layer of the heterogeneous database. It focuses on how to make full and reasonable use of the computing and storage capabilities of existing databases rather than creating a brand new database. It attaches greater importance to the collaboration between multiple databases instead of the database itself.

ShardingSphere-JDBC

ShardingSphere-JDBC is a lightweight Java framework that provides additional services at Java's JDBC layer.

Download PDF



- What is ShardingSphere
 - Introduction
 - Product Features
 - Advantages
 - Roadmap
 - How to Contribute
- Design Philosophy
 - Connect: Create database upper level standard
 - Enhance: Database computing enhancement engine
 - Pluggable: Building database function ecology
- Deployment
 - Deployment
 - Running Modes

ShardingSphere clients


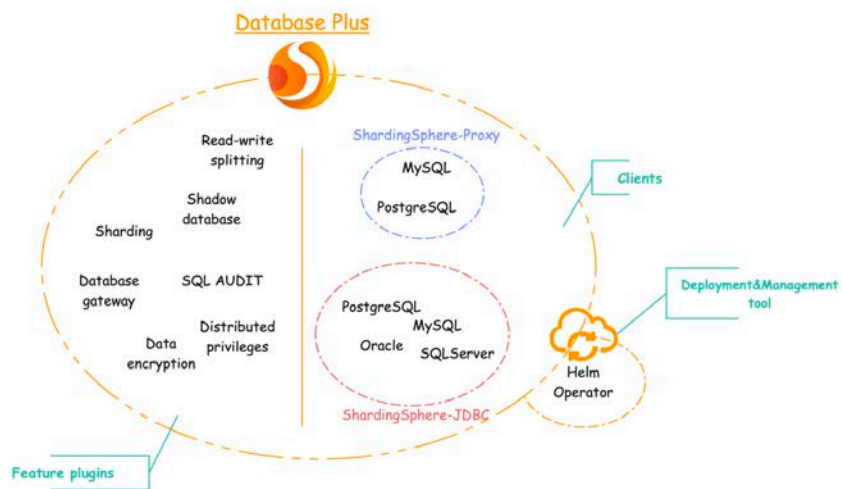
Database Plus

What is Apache ShardingSphere?

The ecosystem to transform any database into a distributed database system, and enhance it with sharding, elastic scaling, encryption features & more.



ShardingSphere features




ShardingSphere on Cloud

ShardingSphere-on-Cloud



Take Apache ShardingSphere to the cloud

A collection of tools & best practices including automated deployment scripts to virtual machines in AWS, Google Cloud Platform, Alibaba Cloud, CloudFormation Stack templates, and Terraform one-click deployment scripts.

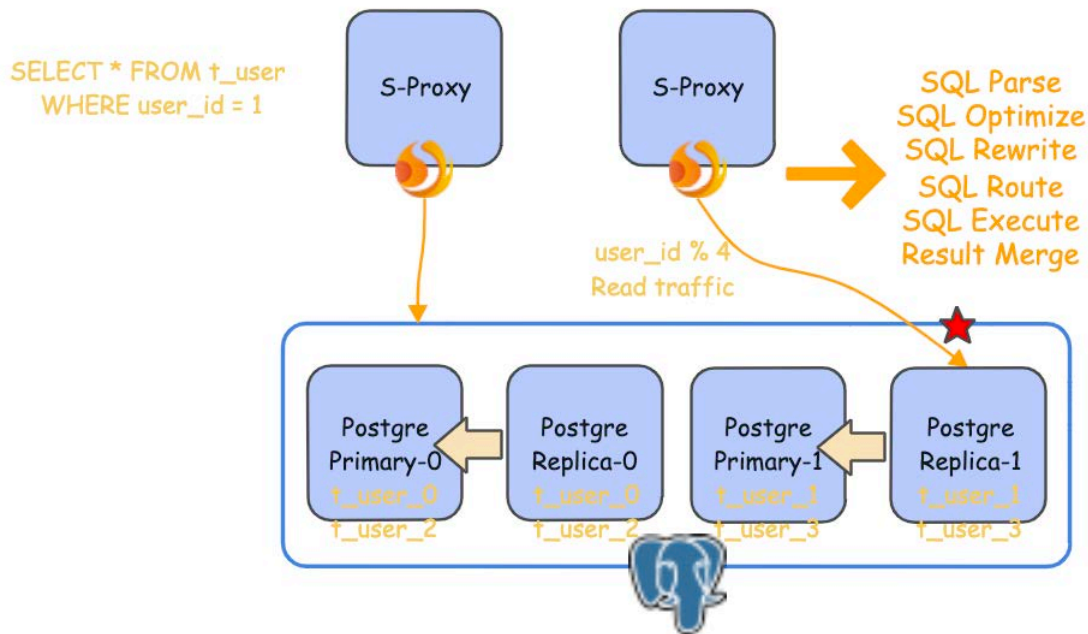
Helm Charts, Operators, automatic horizontal scaling, and other tools for the Kubernetes cloud-native environment are also included.



Demo



The handling process of one SQL



The demo show

1. Deploy two PostgreSQL (Storage node) clusters made of a primary node and a replica
2. Deploy two ShardingSphere-Proxy (Computing node) and ShardingSphere-governance
3. Add PostgreSQL resources and their relationship into ShardingSphere-Proxy
4. Create sharding table t_user on ShardingSphere-Proxy
5. Show the metadata of this distributed database system
6. INSERT data for test on ShardingSphere-Proxy
7. Preview SELECT routing result
8. Execute SELECT query

Step 1, 2,

git clone <https://github.com/apache/shardingsphere-on-cloud>

cd charts/shardingsphere-operator-cluster

helm dependency build

helm install shardingsphere-cluster shardingsphere-operator-cluster -n sharding-test

Pods 5 items

Name ^	Namespace	Containers	Restarts	Controlled ...
shardingsphere-cluster-apach...	shardingsphere-cluster-apache-shardingsphere-proxy-charts-lrvk5			ReplicaSet
shardingsphere-cluster-apach...	sharding-test	■ ■	0	ReplicaSet
shardingsphere-cluster-zooke...	sharding-test	■	0	StatefulSet
shardingsphere-cluster-zooke...	sharding-test	■	1	StatefulSet
shardingsphere-cluster-zooke...	sharding-test	■	0	StatefulSet

```
trista@Tristas-MacPro ~ % helm install pg-cluster-0 bitnami/postgresql -n ss-new --set global.storageClass=sata-csi-udisk --set replication.n
umSynchronousReplicas=1 --set readReplicas.replicaCount=1 --set architecture=replication

NAME: pg-cluster-0
LAST DEPLOYED: Thu Jul 14 12:13:26 2022
NAMESPACE: ss-new
STATUS: deployed
REVISION: 1
TEST SUITE: None
NOTES:
CHART NAME: postgresql
CHART VERSION: 11.6.16
APP VERSION: 14.4.0
```

Services 3 items

Namespace: sharding-test

Name ^	Namespace	Type	Cluster IP	Ports	External IP	Selector
shardingsphere-c...	sharding-test	ClusterIP	172.17.228.23	3307/TCP	-	app=shar
shardingsphere-c...	sharding-test	ClusterIP	172.17.96.233	2181.client/TCP, ...	-	app.kuber
shardingsphere-c...	sharding-test	ClusterIP	None	2181.client/TCP, ...	-	app.kuber

```
trista@Tristas-MacPro ~ % helm install pg-cluster-1 bitnami/postgresql -n ss-new --set global.storageClass=sata-csi-udisk --set replication.r
umSynchronousReplicas=1 --set readReplicas.replicaCount=1 --set architecture=replication

NAME: pg-cluster-1
LAST DEPLOYED: Thu Jul 14 12:13:40 2022
NAMESPACE: ss-new
STATUS: deployed
REVISION: 1
TEST SUITE: None
NOTES:
CHART NAME: postgresql
CHART VERSION: 11.6.16
APP VERSION: 14.4.0
```

Step 3, 4, 5

```
psql (14.2, server 12.3 SphereEx-DBPlusEngine-Proxy 1.1.0)  
Type "help" for help.
```

```
postgres=> CREATE DATABASE sharding_rw_splitting_db;  
CREATE DATABASE
```

```
postgres=> ADD RESOURCE write_ds_0 (  
    HOST=127.0.0.1,  
    PORT=5430,  
    DB=sharding_rw_splitting_db,  
    USER=postgres,  
    PASSWORD=x0xJ1jSIbN  
) , read_ds_0 (  
    HOST=127.0.0.1,  
    PORT=5431,  
    DB=sharding_rw_splitting_db,  
    USER=postgres,  
    PASSWORD=x0xJ1jSIbN  
) , write_ds_1 (  
    HOST=127.0.0.1,  
    PORT=5432,  
    DB=sharding_rw_splitting_db,  
    USER=postgres,  
    PASSWORD=RHVdPNbsyK  
) , read_ds_1 (  
    HOST=127.0.0.1,  
    PORT=5433,  
    DB=sharding_rw_splitting_db,  
    USER=postgres,  
    PASSWORD=RHVdPNbsyK  
);  
SUCCESS
```

```
postgres=>  
postgres=> CREATE READWRITE_SPLITTING RULE rw_group_0 (  
    WRITE_RESOURCE=write_ds_0,  
    READ_RESOURCES(read_ds_0),  
    TYPE(NAME=random)  
);  
SUCCESS
```

```
postgres=> CREATE READWRITE_SPLITTING RULE rw_group_1 (  
    WRITE_RESOURCE=write_ds_1,  
    READ_RESOURCES(read_ds_1),  
    TYPE(NAME=random)  
);  
SUCCESS
```

```
sharding_rw_splitting_db=> CREATE SHARDING TABLE RULE t_user (  
    RESOURCES(rw_group_0,rw_group_1),  
    SHARDING_COLUMN=user_id,TYPE(NAME=mod,PROPERTIES("sharding-count"=4)))  
);  
SUCCESS
```

```
postgres=>  
postgres=> CREATE TABLE t_user (  
    user_id int4,  
    user_name varchar(32),  
    tel varchar(32)  
);  
CREATE TABLE  
postgres=>
```

```
sharding_rw_splitting_db=> SHOW SHARDING TABLE NODES;  
name | nodes  
-----+-----  
t_user | rw_group_0.t_user_0, rw_group_1.t_user_1, rw_group_0.t_user_2, rw_group_1.t_user_3  
(1 row)
```

Step 6, 7, 8

```
postgres=>
postgres=> INSERT INTO t_user values (1,'name1','tel11111');
INSERT INTO t_user values (2,'name2','tel22222');
INSERT INTO t_user values (3,'name3','tel33333');
INSERT INTO t_user values (4,'name4','tel44444');
INSERT 0 1
INSERT 0 1
INSERT 0 1
INSERT 0 1
```

```
sharding_rw_splitting_db=> PREVIEW SELECT * FROM t_user WHERE user_id=1;
data_source_name |          actual_sql
-----+-----
read_ds_1        | SELECT * FROM t_user_1 WHERE user_id=1
(1 row)
```

```
sharding_rw_splitting_db=>
sharding_rw_splitting_db=> SELECT * FROM t_user WHERE user_id=1;
user_id | user_name | tel
-----+-----+-----
1 | name1     | tel11111
(1 row)
```

```
sharding_rw_splitting_db=>
sharding_rw_splitting_db=> PREVIEW SELECT * FROM t_user;
data_source_name |          actual_sql
-----+-----
read_ds_0        | SELECT * FROM t_user_0 UNION ALL SELECT * FROM t_user_2
read_ds_1        | SELECT * FROM t_user_1 UNION ALL SELECT * FROM t_user_3
(2 rows)
```

```
sharding_rw_splitting_db=> SELECT * FROM t_user ORDER BY user_id;
user_id | user_name | tel
-----+-----+-----
1 | name1     | tel11111
2 | name2     | tel22222
3 | name3     | tel33333
4 | name4     | tel44444
(4 rows)

sharding_rw_splitting_db=>
```

Thanks!

Any questions?

Bio: <https://tristazero.github.io>

LinkedIn: <https://www.linkedin.com/in/panjuan>

GitHub: <https://github.com/tristaZero>

Twitter: @tristaZero

Project Twitter: @ShardingSphere