



Low Overhead Python Application Profiling Using eBPF

Yonatan Goldschmidt

Principal Engineer and Research Team Lead @ Granulate

Connect on [LinkedIn](#) | GitHub: <https://github.com/Jongy>

Meet Yonatan Goldschmidt

- About me:

- Six years of service as R&D specialist & team lead
- Likes everything about computers and software
- Today, a team lead on Granulate's Performance Research Team



- About Granulate:

Enables companies to optimize their workloads, improve performance and leverage that to reduce compute costs

Me in Italy, during a sommelier course!

Why profiling is amazing

It's a win-win... win.

- Easy to apply and use
- Gain visibility on code hot spots, bottlenecks
- Identify impactful performance improvement opportunities

Deterministic Python profilers


For versatile profiling needs

- **cProfile:** Deterministic profiling of Python programs, generate Python execution reports in the CLI via the pstats module, or use visualization tools like snakeviz
- Many other Python profilers: Pyinstrument, Yappi and more.

409515 function calls (405541 primitive calls) in 3.421 seconds

Ordered by: internal time

<u>ncalls</u>	<u>tottime</u>	<u>percall</u>	<u>cumtime</u>	<u>percall</u>	<u>filename:lineno(function)</u>
15	0.397	0.026	0.399	0.027	search_index.py:76(<u>not_operation</u>)
1	0.392	0.392	1.186	1.186	/System/Library/Frameworks/Python
1	0.228	0.228	0.532	0.532	dictionary.py:53(<u>from_json</u>)
1	0.193	0.193	0.193	0.193	/System/Library/Frameworks/Python
1949	0.169	0.000	0.169	0.000	{method 'seek' of 'file' objects}
1	0.129	0.129	0.129	0.129	/Library/Python/2.7/site-packages
1	0.116	0.116	0.119	0.119	/System/Library/Frameworks/Python
1	0.110	0.110	0.112	0.112	/System/Library/Frameworks/Python
1	0.085	0.085	0.291	0.291	/System/Library/Frameworks/Python
35735	0.074	0.000	0.074	0.000	dictionary.py:67(<u>_init_</u>)
1	0.060	0.060	0.060	0.060	/System/Library/Frameworks/Python
1	0.059	0.059	1.209	1.209	search_index.py:9(<u>search</u>)
1	0.059	0.059	0.089	0.089	/Library/Python/2.7/site-packages
1	0.048	0.048	0.050	0.050	/System/Library/Frameworks/Python
1	0.045	0.045	0.079	0.079	/Library/Python/2.7/site-packages
1	0.040	0.040	0.219	0.219	/System/Library/Frameworks/Python
1	0.038	0.038	0.040	0.040	/Library/Python/2.7/site-packages
1	0.037	0.037	2.173	2.173	/Library/Python/2.7/site-packages



Deterministic profiler limitations

- Intrusive by design
- Require code changes in many cases (or deployment changes)
- Prone to high overhead (2x slower code)
- Not for production use!

Statistical Python profilers

Lower overhead, and not necessarily intrusive

- **py-spy**: Sampling profiler for Python.
Not intrusive and incurs zero overhead on the application (but has **overhead** on the system). Production safe.
- Other options: vmprof, scalene
(**intrusive**)



Statistical profiler limitations

- Can be made non-intrusive
 - No code changes!
- Low overhead on the system and negligible overhead on the application (if any)
- Suitable for production use
- When is deterministic profiling preferred?



**eBPF changed the profiling
game**

About eBPF

- **1992:** Berkeley Packet Filter released, obscure kernel technology
- **2022:** Now known as eBPF and a household name, in-kernel execution environment:
 - User-defined programs
 - Limited, secure kernel access

```
# tcpdump -d host 127.0.0.1 and port 80
(000) ldh      [12]
(001) jeq      #0x800          jt 2   jf 18
(002) ld       [26]
(003) jeq      #0x7f000001     jt 6   jf 4
(004) ld       [30]
(005) jeq      #0x7f000001     jt 6   jf 18
(006) ldb      [23]
(007) jeq      #0x84          jt 10  jf 8
(008) jeq      #0x6           jt 10  jf 9
(009) jeq      #0x11          jt 10  jf 18
(010) ldh      [20]
(011) jset     #0x1fff         jt 18  jf 12
(012) ldx      4*([14]&0xf)
(013) ldh      [x + 14]
(014) jeq      #0x50          jt 17  jf 15
(015) ldh      [x + 16]
(016) jeq      #0x50          jt 17  jf 18
(017) ret      #262144
(018) ret      #0
```

https://www.usenix.org/system/files/lisa21_slides_gregg_bpf.pdf

eBPF – low overhead, system-wide tracing

PID	COMM	FD	ERR	PATH
490346	ThreadPoolForeg	118	0	/home/jong/.cache/pgadmin4/Default/Cache/df089dba81afe576_0
490396	python3	11	0	/home/jong/.pgadmin/pgadmin4.db
5581	Chrome_IOThread	175	0	/dev/shm/.com.google.Chrome.Quillb
5581	Chrome_IOThread	186	0	/dev/shm/.com.google.Chrome.Quillb
490396	python3	14	0	/home/jong/.pgadmin/sessions/6be80f34-9d9c-4268-8877-2277d86ecc05
490312	Chrome_IOThread	224	0	/dev/shm/.io.nwjs.GNJ0Vj
490312	Chrome_IOThread	226	0	/dev/shm/.io.nwjs.GNJ0Vj
490346	ThreadPoolForeg	118	0	/home/jong/.config/pgadmin4/Default/Cookies-journal
490346	ThreadPoolForeg	120	0	/home/jong/.config/pgadmin4/Default
490346	ThreadPoolForeg	118	0	/home/jong/.cache/pgadmin4/Default/Cache/df089dba81afe576_0
490396	python3	11	0	/home/jong/.pgadmin/pgadmin4.db
490346	ThreadPoolForeg	123	0	/home/jong/.cache/pgadmin4/Default/Cache/06a5699cbd40e70f_0
490396	python3	15	0	/home/jong/.pgadmin/pgadmin4.db
490363	exe	132	0	/home/jong/.local/share/pgadmin/pgadmin4.1642091395220.log
490396	python3	16	0	/home/jong/.pgadmin/sessions/6be80f34-9d9c-4268-8877-2277d86ecc05
490312	Chrome_IOThread	224	0	/dev/shm/.io.nwjs.3srpzk
490312	Chrome_IOThread	226	0	/dev/shm/.io.nwjs.3srpzk
490312	Chrome_IOThread	224	0	/dev/shm/.io.nwjs.L7utri
490312	Chrome_IOThread	226	0	/dev/shm/.io.nwjs.L7utri

```
18 BEGIN
19 {
20     printf("Tracing open syscalls... Hit Ctrl-C to end.\n");
21     printf("%-6s %-16s %4s %3s %s\n", "PID", "COMM", "FD", "ERR", "PATH");
22 }
23
24 tracepoint:syscalls:sys_enter_open,
25 tracepoint:syscalls:sys_enter_openat
26 {
27     @filename[tid] = args->filename;
28 }
29
30 tracepoint:syscalls:sys_exit_open,
31 tracepoint:syscalls:sys_exit_openat
32 /@filename[tid]/
33 {
34     $ret = args->ret;
35     $fd = $ret > 0 ? $ret : -1;
36     $errno = $ret > 0 ? 0 : - $ret;
37
38     printf("%-6d %-16s %4d %3d %s\n", pid, comm, $fd, $errno,
39         str(@filename[tid]));
40     delete(@filename[tid]);
41 }
42
43 END
44 {
45     clear(@filename);
46 }
```

eBPF: A new type of software

	Execution model	User defined	Compilation	Security	Failure mode	Resource access
User	task	yes	any	user based	abort	syscall, fault
Kernel	task	no	static	none	panic	direct
BPF	event	yes	JIT, CO-RE	verified, JIT	error message	restricted helpers

Python profiling in eBPF

- Let's use eBPF to write a better profiler!
- The story of PyPerf – an open-source profiling tool for Python
 - py-spy is good, but we needed something more capable
 - Found PoC profilers PyPerf and rbperf
 - Took PyPerf miles ahead by adding many new features

picsapp-2-optimized

Profile Type

Process names

Weight

2021-08-18

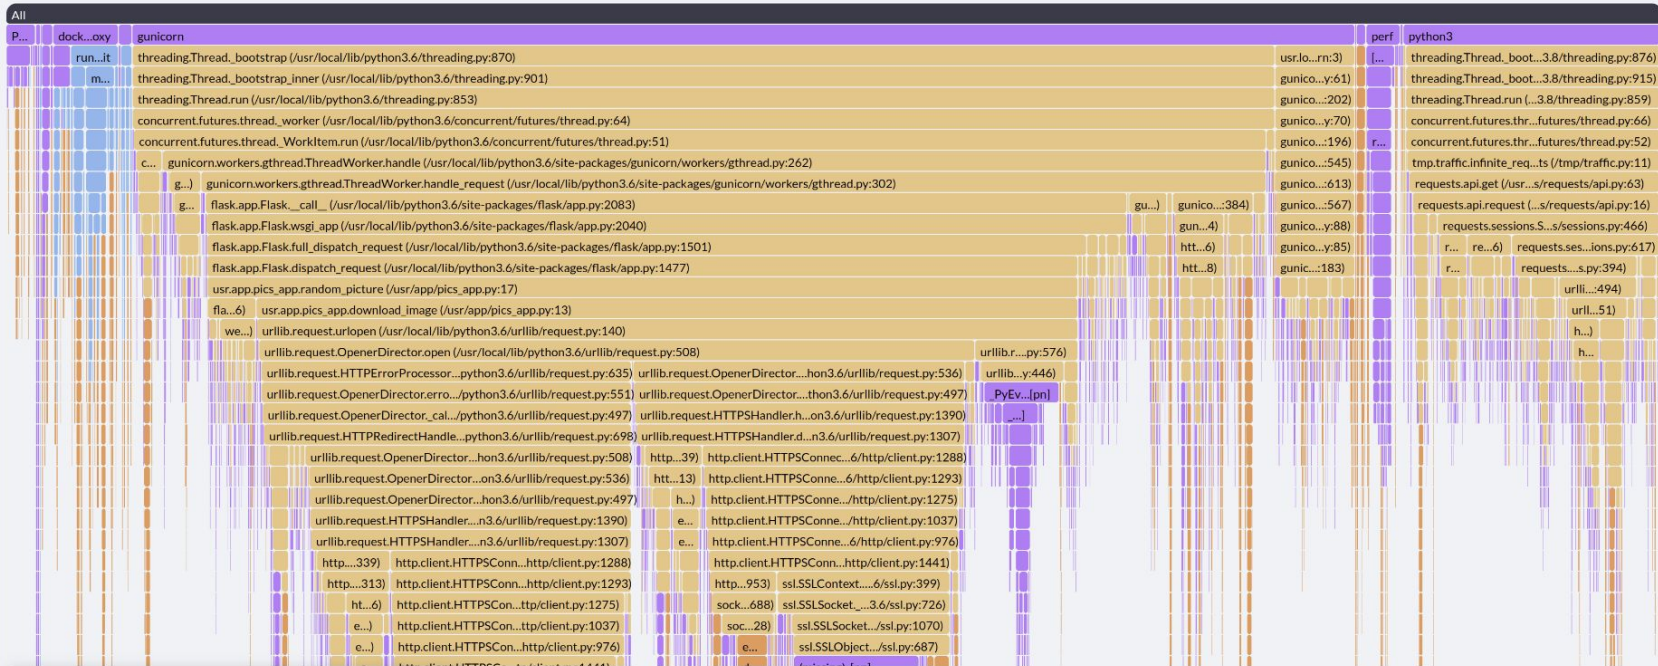


Search

Actions

Name root | samples (0% out of 6068 in total)

Reset view



Java Python PHP C++ Kernel Go Other

Native code matters!

```
import zlib

def func(n):
    buf = ""
    for i in range(n):
        buf += chr(i % 128)
    zlib.compress(buf.encode())

while True: func(999999)
```


Native code matters!

Observing the native side of Python
helps improve performance in
unforeseen ways

% operator


chr()

compression


Loading
global
variables,
e.g "chr"

str +=

1. str +=
2. chr()
3. %
4. compress



PyPerf is a part of gProfiler:



a free & open-source profiler

Thank You!

Feel free to DM here or connect on [LinkedIn](#) and [GitHub](#) (/jongy)

To start free profiling, visit: gprofiler.io





Q&A